

Exercise 2

1. *SYMBOL RESOLUTION*. The following program consists of two modules: `foo` and `bar`. The source code are shown below.

```

/** foo.c */
#include <stdio.h>
void f(void);
short a = 0x1;
short b;
static short c = 0x3;
int main(void) {
    b = 0x2;
    short d = 0x4;
    static int e = 0x10;
    f();
    printf("a=0x%x b=0x%x c=0x%x d=0x%x e=0x%x\n", a, b, c, d, e);
    return 0;
}

/** bar.c */
long a;
int d;
void f(void) {
    a = 0x0;
    d = 0x0;
    int e = 0x0;
}

```

- (a) For each symbol in `foo.o`, please indicate whether it will have a symbol table entry in the `.symtab` section. If Yes, please fill the binding (GLOBAL, LOCAL); If No, fill with ‘-’.

Sym Name	Has a .symtab Entry?	Binding
a	<u>Y</u>	GLOBAL
b	<u>Y</u>	GLOBAL
c	<u>Y</u>	LOCAL
d	<u>N</u>	--
e	<u>Y</u>	LOCAL
f	<u>Y</u>	GLOBAL

- (b) `foo.o` and `bar.o` are linked to `foobar`. The output of `readelf -s foobar` is provided below (Some entries are omitted). What is the output after running `./foobar`?

Num:	Value	Size	Type	Bind	Vis	Ndx	Name
142:	000000000006b90f2	2	OBJECT	LOCAL	DEFAULT	21	c
143:	000000000006b90f4	4	OBJECT	LOCAL	DEFAULT	21	e.2256
746:	000000000006bc3a0	2	OBJECT	GLOBAL	DEFAULT	26	b
834:	00000000000400baf	35	FUNC	GLOBAL	DEFAULT	6	f
1376:	00000000000400b4d	98	FUNC	GLOBAL	DEFAULT	6	main
1425:	000000000006bc3a4	4	OBJECT	GLOBAL	DEFAULT	26	d
1633:	000000000006b90f0	2	OBJECT	GLOBAL	DEFAULT	21	a

a=0x0 b=0x2 c=0x0 d=0x4 e=0x0

2. *ELF INSIDE*. For a source file `a.c`, an ELF object file `a.o` is derived using `gcc -c a.c` (compiled and assembled, but not linked). Here is the source code and a disassembly of the `.text` section of `a.o`.

```

long seq[3] = {1, 2, 3};
long flag;

int main(int argc, char **argv)
{
    if (seq[2] > 0) {
        flag = 1;
    }
    else {
        flag = 0;
    }
}

```

```

    }
    return 0;
}

Disassembly of section .text:
0000000000000000 <main>:
 0: 55          push    %rbp
 1: 48 89 e5    mov     %rsp,%rbp
 4: 89 7d fc    mov     %edi,-0x4(%rbp)
 7: 48 89 75 f0  mov     %rsi,-0x10(%rbp)
 b: 48 8b 05 00 00 00 00 mov     0x0(%rip),%rax
12: 48 85 c0    test    %rax,%rax
15: 7e 0d      jle     24 <main+0x24>
17: 48 c7 05 00 00 00 00 movq    $0x1,0x0(%rip)
1e: 01 00 00 00
22: eb 0b      jmp     2f <main+0x2f>
24: 48 c7 05 00 00 00 00 movq    $0x0,0x0(%rip)
2b: 00 00 00 00
2f: b8 00 00 00 00 mov     $0x0,%eax
34: 5d        pop     %rbp
35: c3        retq

```

- (a) What are the symbol table entries for symbol `seq`, `flag` and `main`?

Sym Name	Section/Pseudosection	Type	Binding	Size
<code>seq</code>	<code>.data</code>	<code>OBJ</code>	<code>GLOBAL</code>	<code>24</code>
<code>flag</code>	<code>.COM</code>	<code>OBJ</code>	<code>GLOBAL</code>	<code>8</code>
<code>main</code>	<code>.text</code>	<code>FUNC</code>	<code>GLOBAL</code>	<code>54</code>

- (b) What is the relocation entry for the position at `0xe` after the `.text` section?

Offset	Sym	Type	Addend
<code>0xe</code>	<code>seq</code>	<code>R_X86_64_PC32</code>	<code>0xc</code>

- (c) When linking, suppose the linker has decided that the address of `.text` is `0x400b4d`, the address of `seq` is `0x6b90f0`. What is the relocated form of the `mov 0x0(%rip),%rax` instruction?

$0x6b90f0 - (0x400b4d + 0xe) + 0xc = 0x2b85a1$. So in the resulting executable object file, the relocated form of this instruction is `48 8b 05 a1 85 2b 00`.

- (d) Clark prepares to do some hacking to the ELF file `a.o`. He wants to manually manipulate the `Type` field of the relocation entry in question (b) from `R_X86_64_PC32` to `R_X86_64_32`. Please try your best to give a guess of how he does it step-by-step. You don't need to worry about the details of file-related operations. For your information, it first opens `a.o`, then after `mmap`, the file data is "mapped" into the memory. Memory accesses to this region are all backed by the underlying file.

```

if ((fd = open(file_name, O_RDWR)) < 0) {
    fprintf(stderr, "Cannot open file %s\n", file_name);
    return errno;
}

fstat(fd, &sb);
file_mapped = mmap(NULL, sb.st_size, PROT_READ | PROT_WRITE,
    MAP_SHARED, fd, 0);
if (file_mapped == MAP_FAILED) {
    fprintf(stderr, "mmap\n");
    return errno;
}

ehdr = (Elf64_Ehdr *)file_mapped;
shdrs = (Elf64_Shdr *) (file_mapped + ehdr->e_shoff);
for (i = 0; i < ehdr->e_shnum; i++) {
    if (shdrs[i].sh_type == SHT_RELA)
        break;
}

```

```

if (i >= ehdr->e_shnum) {
    fprintf(stderr, "rel not found\n");
    return -1;
}

rel = (Rel_ent *) (file_mapped + shdrs[i].sh_offset);
rel[0].r_type = R_X86_64_32;

munmap(file_mapped, sb.st_size);
close(fd);

```

- `ehdr` represents the ELF header which is at the beginning of the ELF file.
- In ELF header, there is a field (`ehdr->e_shoff`) tells the byte offset of the section header table.
- The section header table is an array of section headers (`shdrs`).
- Each section header records the information about one section. The field about section type is searched for the relocation section (`SHT_RELA`).
- Once found, use the offset in the section header (`shdrs[i].sh_offset`) to locate the relocation section.
- The relocation section is an array of reloc entries. Change the first entry's type field to `R_X86_64_32`.

(e) After (d), what is the relocated form of the `mov 0x0(%rip),%rax` instruction? (Suppose the address of `.text` and `seq` is the same as question (c))

$0x6b90f0 + 0xc = 0x6b90fc$. So in the resulting executable object file, the relocated form of this instruction is `48 8b 05 fc 90 6b 00`.

3. **RELOCATION.** Two source files and the disassembly of their object files are given below.

```

/** foo.c **/
static int n = 2013;
int *p_n = &n;
int foo(int x) {
    if (x < n) return 1;
    return foo(x-1) * n;
}

/** bar.c **/
extern int foo(int n);
extern int *p_n;
int n = 2015;
int a[2048];
void bar(void) {
    *p_n = 2014;
    a[2] = foo(n);
}

/** foo.obj **/
0000000000000000 <foo>:
0: 55                push    %rbp
1: 48 89 e5          mov     %rsp,%rbp
4: 48 83 ec 10       sub     $0x10,%rsp
8: 89 7d fc          mov     %edi,-0x4(%rbp)
b: 8b 05 00 00 00 00 mov     0x0(%rip),%eax    @@
11: 39 45 fc          cmp     %eax,-0x4(%rbp)
14: 7d 07            jge     1d <foo+0x1d>
16: b8 01 00 00 00    mov     $0x1,%eax
1b: eb 18            jmp     35 <foo+0x35>
1d: 8b 45 fc          mov     -0x4(%rbp),%eax
20: 83 e8 01          sub     $0x1,%eax
23: 89 c7            mov     %eax,%edi

```

25:	e8 00 00 00 00	callq	2a <foo+0x2a>	@@
2a:	89 c2	mov	%eax,%edx	
2c:	8b 05 00 00 00 00	mov	0x0(%rip),%eax	
32:	0f af c2	imul	%edx,%eax	
35:	c9	leaveq		
36:	c3	retq		
/** bar.obj **/				
0000000000000000 <bar>:				
0:	55	push	%rbp	
1:	48 89 e5	mov	%rsp,%rbp	
4:	48 8b 05 00 00 00 00	mov	0x0(%rip),%rax	@@
b:	c7 00 de 07 00 00	movl	\$0x7de, (%rax)	
11:	8b 05 00 00 00 00	mov	0x0(%rip),%eax	
17:	89 c7	mov	%eax,%edi	
19:	e8 00 00 00 00	callq	1e <bar+0x1e>	
1e:	89 05 00 00 00 00	mov	%eax,0x0(%rip)	@@
24:	90	nop		
25:	5d	pop	%rbp	
26:	c3	retq		

(a) Fill in the symbol table of foo.o.

Type	Binding	Section/Pseudosection	Name
OBJ	<u>LOCAL</u>	.data	n
OBJ	GLOBAL	<u>.data</u>	p_n
FUNC	<u>GLOBAL</u>	<u>.text</u>	foo

(b) Fill in the symbol table of bar.o.

Type	Binding	Section/Pseudosection	Name
OBJ	<u>GLOBAL</u>	.data	n
OBJ	GLOBAL	<u>COMM</u>	a
OBJ	<u>GLOBAL</u>	<u>UNDEF</u>	p_n
FUNC	<u>GLOBAL</u>	<u>UNDEF</u>	foo

(c) Fill in the relocation entries of foo.o.

Section	Offset	Type	Addend
.text	<u>0x0000000d</u>	<u>R_x86_64_PC32</u>	-4
.text	0x00000026	R_x86_64_PC32	<u>-4</u>
.text	0x0000002e	<u>R_x86_64_PC32</u>	<u>-4</u>

(d) Fill in the relocation entries of bar.o.

Section	Offset	Type	Addend
.text	0x00000007	<u>R_x86_64_PC32</u>	-4
.text	0x00000013	<u>R_x86_64_PC32</u>	-4
.text	<u>0x0000001a</u>	<u>R_x86_64_PC32</u>	<u>-4</u>
.text	<u>0x00000020</u>	R_x86_64_PC32	<u>+4</u>

(e) After relocation and the program is built, what is the relocated form of the 4 instructions tagged with @@? Suppose the runtime address of some symbols are decided as below.

foo	0x4004fd
bar	0x4004d6
n (foo.o's .data)	0x601038
n (bar.o's .data)	0x601088
p_n	0x601040
a	0x601080

1. 2a 0b 20 00 (0x601038 - 0x4004fd - 0x11 = 0x200b2a)
2. d6 ff ff ff (0x4004fd - (0x4004fd + 0x26) - 0x4 = 0xffffffff5)
3. 5f 0b 20 00 (0x601040 - 0x4004dd - 4 = 0x200b5f)
4. 8e 0b 20 00 (0x601080 - 0x4004f6 + 4 = 0x200b8e)