

Exercise 2

1. *SYMBOL RESOLUTION*. The following program consists of two modules: `foo` and `bar`. The source code are shown below.

```
/** foo.c */
#include <stdio.h>
void f(void);
short a = 0x1;
short b;
static short c = 0x3;
int main(void) {
    b = 0x2;
    short d = 0x4;
    static int e = 0x10;
    f();
    printf("a=0x%x b=0x%x c=0x%x d=0x%x e=0x%x\n", a, b, c, d, e);
    return 0;
}

/** bar.c */
long a;
int d;
void f(void) {
    a = 0x0;
    d = 0x0;
    int e = 0x0;
}
```

- (a) For each symbol in `foo.o`, please indicate whether it will have a symbol table entry in the `.symtab` section. If Yes, please fill the binding (GLOBAL, LOCAL); If No, fill with ‘-’.

Sym Name	Has a .symtab Entry?	Binding
a	_____	_____
b	_____	_____
c	_____	_____
d	_____	_____
e	_____	_____
f	_____	_____

- (b) `foo.o` and `bar.o` are linked to `foobar`. The output of `readelf -s foobar` is provided below (Some entries are omitted). What is the output after running `./foobar`?

Num:	Value	Size	Type	Bind	Vis	Ndx	Name
142:	000000000006b90f2	2	OBJECT	LOCAL	DEFAULT	21	c
143:	000000000006b90f4	4	OBJECT	LOCAL	DEFAULT	21	e.2256
746:	000000000006bc3a0	2	OBJECT	GLOBAL	DEFAULT	26	b
834:	00000000000400baf	35	FUNC	GLOBAL	DEFAULT	6	f
1376:	00000000000400b4d	98	FUNC	GLOBAL	DEFAULT	6	main
1425:	000000000006bc3a4	4	OBJECT	GLOBAL	DEFAULT	26	d
1633:	000000000006b90f0	2	OBJECT	GLOBAL	DEFAULT	21	a

2. *ELF INSIDE*. For a source file `a.c`, an ELF object file `a.o` is derived using `gcc -c a.c` (compiled and assembled, but not linked). Here is the source code and a disassembly of the `.text` section of `a.o`.

```
long seq[3] = {1, 2, 3};
long flag;

int main(int argc, char **argv)
{
    if (seq[2] > 0) {
        flag = 1;
    }
    else {
        flag = 0;
    }
}
```

```

    }
    return 0;
}

Disassembly of section .text:
0000000000000000 <main>:
 0: 55          push    %rbp
 1: 48 89 e5    mov     %rsp,%rbp
 4: 89 7d fc    mov     %edi,-0x4(%rbp)
 7: 48 89 75 f0  mov     %rsi,-0x10(%rbp)
 b: 48 8b 05 00 00 00 00 mov     0x0(%rip),%rax
12: 48 85 c0    test    %rax,%rax
15: 7e 0d      jle     24 <main+0x24>
17: 48 c7 05 00 00 00 00 movq    $0x1,0x0(%rip)
1e: 01 00 00 00
22: eb 0b      jmp     2f <main+0x2f>
24: 48 c7 05 00 00 00 00 movq    $0x0,0x0(%rip)
2b: 00 00 00 00
2f: b8 00 00 00 00 mov     $0x0,%eax
34: 5d        pop     %rbp
35: c3        retq

```

- (a) What are the symbol table entries for symbol `seq`, `flag` and `main`?

Sym Name	Section/Pseudosection	Type	Binding	Size
<code>seq</code>	_____	_____	_____	_____
<code>flag</code>	_____	_____	_____	_____
<code>main</code>	_____	_____	_____	54

- (b) What is the relocation entry for the position at `0xe` after the `.text` section?

Offset	Sym	Type	Addend
<code>0xe</code>	<code>seq</code>	<code>R_X86_64_PC32</code>	_____

- (c) When linking, suppose the linker has decided that the address of `.text` is `0x400b4d`, the address of `seq` is `0x6b90f0`. What is the relocated form of the `mov 0x0(%rip),%rax` instruction?
- (d) Clark prepares to do some hacking to the ELF file `a.o`. He wants to manually manipulate the `Type` field of the relocation entry in question (b) from `R_X86_64_PC32` to `R_X86_64_32`. Please try your best to give a guess of how he does it step-by-step. You don't need to worry about the details of file-related operations. For your information, it first opens `a.o`, then after `mmap`, the file data is "mapped" into the memory. Memory accesses to this region are all backed by the underlying file.

```

if ((fd = open(file_name, O_RDWR)) < 0) {
    fprintf(stderr, "Cannot open file %s\n", file_name);
    return errno;
}

fstat(fd, &sb);
file_mapped = mmap(NULL, sb.st_size, PROT_READ | PROT_WRITE,
    MAP_SHARED, fd, 0);
if (file_mapped == MAP_FAILED) {
    fprintf(stderr, "mmap\n");
    return errno;
}

ehdr = (Elf64_Ehdr *)file_mapped;
shdrs = (Elf64_Shdr *) (file_mapped + ehdr->e_shoff);
for (i = 0; i < ehdr->e_shnum; i++) {
    if (shdrs[i].sh_type == SHT_RELA)
        break;
}

if (i >= ehdr->e_shnum) {

```

```

        fprintf(stderr, "rel not found\n");
        return -1;
    }

    rel = (Rel_ent *) (file_mapped + shdrs[i].sh_offset);

    rel[0].r_type = R_X86_64_32;

    munmap(file_mapped, sb.st_size);
    close(fd);

```

- (e) After (d), what is the relocated form of the `mov 0x0(%rip),%rax` instruction? (Suppose the address of `.text` and `seq` is the same as question (c))

3. *RELOCATION*. Two source files and the disassembly of their object files are given below.

```

/** foo.c */
static int n = 2013;
int *p_n = &n;
int foo(int x) {
    if (x < n) return 1;
    return foo(x-1) * n;
}

/** bar.c */
extern int foo(int n);
extern int *p_n;
int n = 2015;
int a[2048];
void bar(void) {
    *p_n = 2014;
    a[2] = foo(n);
}

/** foo.obj */
0000000000000000 <foo>:
  0: 55                push    %rbp
  1: 48 89 e5          mov     %rsp,%rbp
  4: 48 83 ec 10       sub     $0x10,%rsp
  8: 89 7d fc          mov     %edi,-0x4(%rbp)
 b: 8b 05 00 00 00 00 mov     0x0(%rip),%eax    @@
11: 39 45 fc          cmp     %eax,-0x4(%rbp)
14: 7d 07             jge     1d <foo+0x1d>
16: b8 01 00 00 00    mov     $0x1,%eax
1b: eb 18             jmp     35 <foo+0x35>
1d: 8b 45 fc          mov     -0x4(%rbp),%eax
20: 83 e8 01          sub     $0x1,%eax
23: 89 c7             mov     %eax,%edi
25: e8 00 00 00 00    callq  2a <foo+0x2a>    @@
2a: 89 c2             mov     %eax,%edx
2c: 8b 05 00 00 00 00 mov     0x0(%rip),%eax
32: 0f af c2          imul    %edx,%eax
35: c9               leaveq  %edx,%eax
36: c3               retq

/** bar.obj */
0000000000000000 <bar>:
  0: 55                push    %rbp
  1: 48 89 e5          mov     %rsp,%rbp
  4: 48 8b 05 00 00 00 00 mov     0x0(%rip),%rax    @@
 b: c7 00 de 07 00 00 movl    $0x7de,(%rax)
11: 8b 05 00 00 00 00 mov     0x0(%rip),%eax
17: 89 c7             mov     %eax,%edi
19: e8 00 00 00 00    callq  1e <bar+0x1e>
1e: 89 05 00 00 00 00 mov     %eax,0x0(%rip)    @@
24: 90               nop
25: 5d               pop     %rbp
26: c3               retq

```

(a) Fill in the symbol table of foo.o.

Type	Binding	Section/Pseudosection	Name
OBJ	_____	.data	n
OBJ	GLOBAL	_____	p_n
FUNC	_____	_____	foo

(b) Fill in the symbol table of bar.o.

Type	Binding	Section/Pseudosection	Name
OBJ	_____	.data	n
OBJ	GLOBAL	_____	a
OBJ	_____	_____	p_n
FUNC	_____	_____	foo

(c) Fill in the relocation entries of foo.o.

Section	Offset	Type	Addend
.text	_____	_____	-4
.text	0x00000026	R_x86_64_PC32	_____
.text	0x0000002e	_____	_____

(d) Fill in the relocation entries of bar.o.

Section	Offset	Type	Addend
.text	0x00000007	_____	-4
.text	0x00000013	_____	-4
.text	_____	_____	_____
.text	_____	R_x86_64_PC32	_____

(e) After relocation and the program is built, what is the relocated form of the 4 instructions tagged with @@? Suppose the runtime address of some symbols are decided as below.

foo	0x4004fd
bar	0x4004d6
n (foo.o's .data)	0x601038
n (bar.o's .data)	0x601088
p_n	0x601040
a	0x601080